

house-prediction

August 3, 2025

[]:

1 Task

Develop a regression model to predict house prices using the data from “/content/House Price Prediction Dataset.csv”. The model should be built using Python and should include data loading, preprocessing, exploratory data analysis, model training, and evaluation.

1.1 Data loading and initial exploration

1.1.1 Subtask:

Load the dataset and perform an initial exploration to understand its structure, features, and identify any immediate data quality issues.

Reasoning: I will import the pandas library, load the dataset from the specified path, and then display the first few rows, the info, and the descriptive statistics of the dataframe to complete the initial exploration of the dataset as requested.

[]:

```
import pandas as pd

df = pd.read_csv('/content/House Price Prediction Dataset.csv')
display(df.head())
print(df.info())
display(df.describe())
```

	Id	Area	Bedrooms	Bathrooms	Floors	YearBuilt	Location	Condition	\
0	1	1360	5	4	3	1970	Downtown	Excellent	
1	2	4272	5	4	3	1958	Downtown	Excellent	
2	3	3592	2	2	3	1938	Downtown	Good	
3	4	966	4	2	2	1902	Suburban	Fair	
4	5	4926	1	4	2	1975	Downtown	Fair	

	Garage	Price
0	No	149919
1	No	424998
2	No	266746
3	Yes	244020
4	Yes	636056

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Id           2000 non-null   int64
1   Area         2000 non-null   int64
2   Bedrooms     2000 non-null   int64
3   Bathrooms    2000 non-null   int64
4   Floors       2000 non-null   int64
5   YearBuilt    2000 non-null   int64
6   Location     2000 non-null   object
7   Condition    2000 non-null   object
8   Garage       2000 non-null   object
9   Price        2000 non-null   int64
dtypes: int64(7), object(3)
memory usage: 156.4+ KB
None

```

	Id	Area	Bedrooms	Bathrooms	Floors \
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1000.500000	2786.209500	3.003500	2.55250	1.993500
std	577.494589	1295.146799	1.424606	1.10899	0.809188
min	1.000000	501.000000	1.000000	1.00000	1.000000
25%	500.750000	1653.000000	2.000000	2.00000	1.000000
50%	1000.500000	2833.000000	3.000000	3.00000	2.000000
75%	1500.250000	3887.500000	4.000000	4.00000	3.000000
max	2000.000000	4999.000000	5.000000	4.00000	3.000000

	YearBuilt	Price
count	2000.000000	2000.000000
mean	1961.446000	537676.855000
std	35.926695	276428.845719
min	1900.000000	50005.000000
25%	1930.000000	300098.000000
50%	1961.000000	539254.000000
75%	1993.000000	780086.000000
max	2023.000000	999656.000000

1.2 Data preprocessing and feature engineering

1.2.1 Subtask:

Clean the data by handling missing values and duplicates. Create new features if necessary to improve model performance.

Reasoning: I will start by checking for and removing any duplicate rows in the dataframe. Then, I will convert the categorical features to a numerical format using one-hot encoding, and finally, I will drop the original categorical columns and concatenate the new one-hot encoded columns to

the dataframe. This will complete all the steps of the subtask in a single code block.

```
[ ]: # Check for duplicates
print(f"Number of duplicate rows: {df.duplicated().sum()}")
# Remove duplicates
df.drop_duplicates(inplace=True)
print(f"Number of rows after removing duplicates: {len(df)}")

# One-hot encode categorical features
df_encoded = pd.get_dummies(df, columns=['Location', 'Condition', 'Garage'],
                              drop_first=True)

# Display the first few rows of the new dataframe
display(df_encoded.head())
```

Number of duplicate rows: 0

Number of rows after removing duplicates: 2000

	Id	Area	Bedrooms	Bathrooms	Floors	YearBuilt	Price	Location_Rural	\
0	1	1360	5	4	3	1970	149919	False	
1	2	4272	5	4	3	1958	424998	False	
2	3	3592	2	2	3	1938	266746	False	
3	4	966	4	2	2	1902	244020	False	
4	5	4926	1	4	2	1975	636056	False	

	Location_Suburban	Location_Urban	Condition_Fair	Condition_Good	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	True	
3	True	False	True	False	
4	False	False	True	False	

	Condition_Poor	Garage_Yes
0	False	False
1	False	False
2	False	False
3	False	True
4	False	True

1.3 Exploratory data analysis (eda)

1.3.1 Subtask:

Conduct a more in-depth analysis of the data to uncover relationships between different features and the target variable (price). This will involve creating visualizations and summary statistics.

Reasoning: Calculate and display the correlation matrix for `df_encoded` to understand the linear relationships between features, including the target variable 'Price', and then create various plots to visualize the relationships between features and the target variable.

```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns



# Calculate and display the correlation matrix
correlation_matrix = df_encoded.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Features and Price')
plt.show()

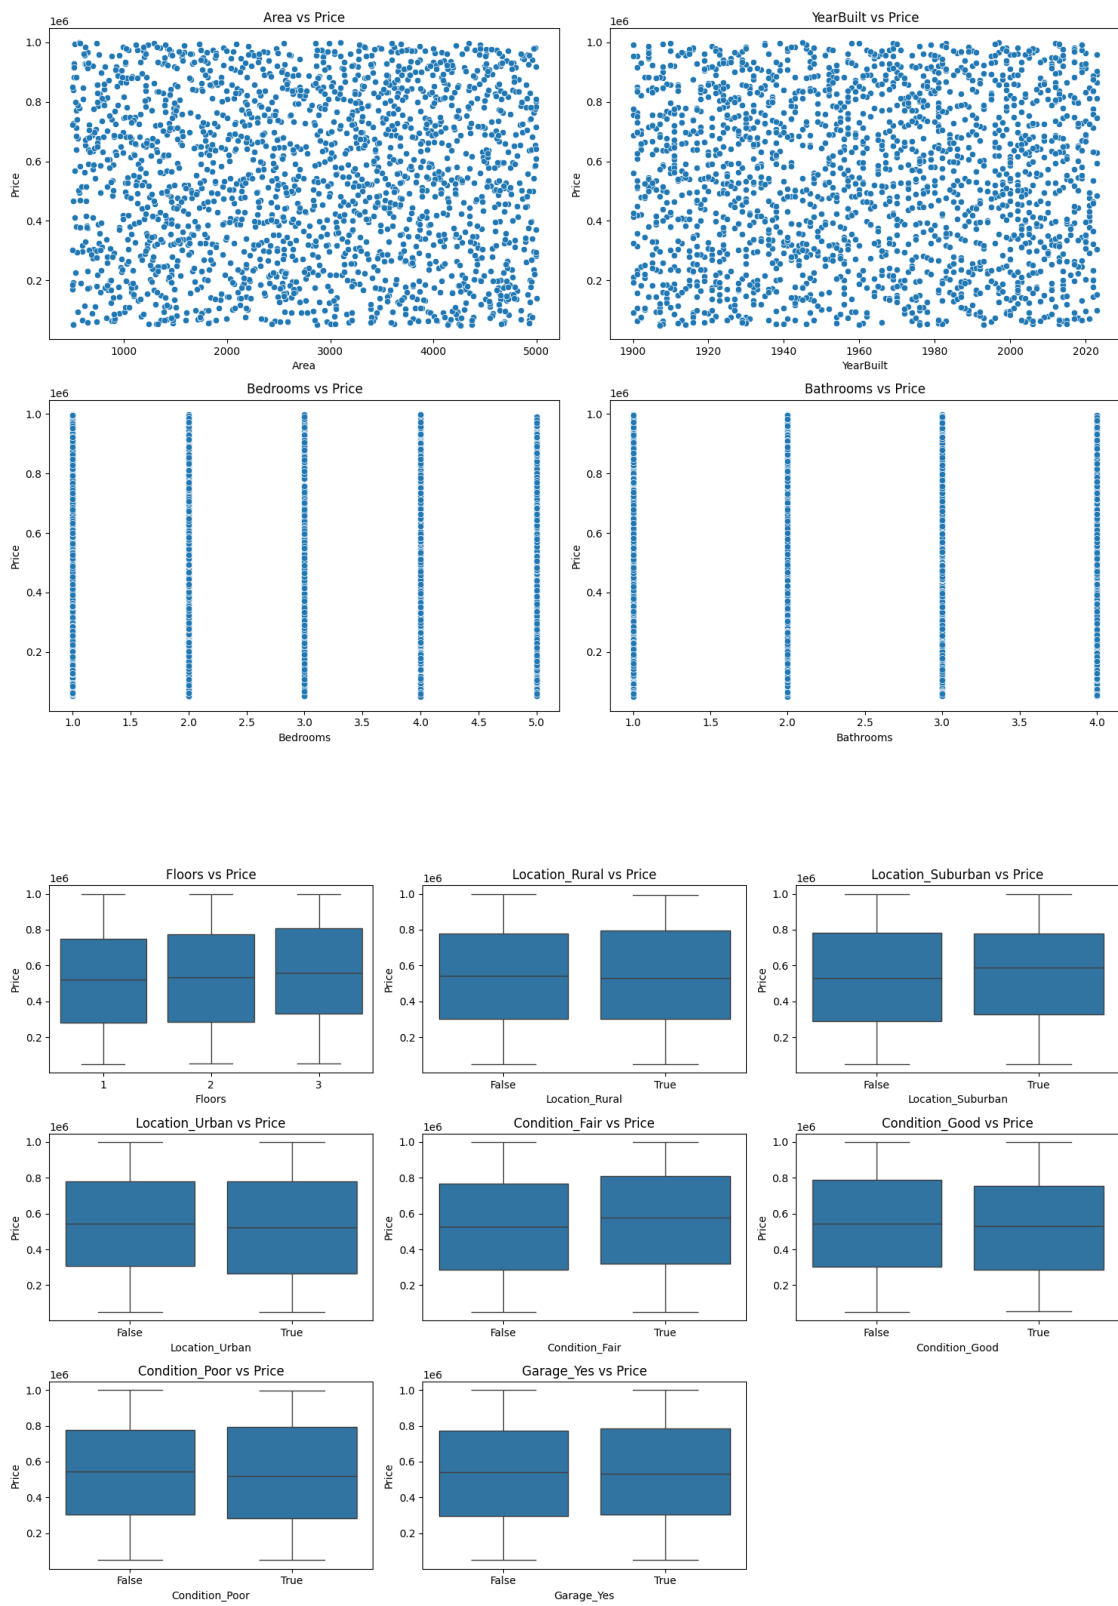
# Create a histogram of the 'Price' column
plt.figure(figsize=(10, 6))
sns.histplot(df_encoded['Price'], kde=True)
plt.title('Distribution of House Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()

# Generate scatter plots for numerical features against 'Price'
numerical_features = ['Area', 'YearBuilt', 'Bedrooms', 'Bathrooms']
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numerical_features):
    plt.subplot(2, 2, i + 1)
    sns.scatterplot(x=df_encoded[feature], y=df_encoded['Price'])
    plt.title(f'{feature} vs Price')
plt.tight_layout()
plt.show()

# Create box plots for one-hot encoded categorical features against 'Price'
categorical_features_encoded = [col for col in df_encoded.columns if col not in
    numerical_features + ['Id', 'Price']]

plt.figure(figsize=(15, 10))
for i, feature in enumerate(categorical_features_encoded):
    plt.subplot(3, 3, i + 1)
    sns.boxplot(x=df_encoded[feature], y=df_encoded['Price'])
    plt.title(f'{feature} vs Price')
plt.tight_layout()
plt.show()
```





1.4 Model training

1.4.1 Subtask:

Split the data into training and testing sets. Train a regression model on the training data. I will start with a simple linear regression model.

Reasoning: Import the necessary modules for splitting the data and training a linear regression model, define features and the target, split the data, instantiate the model, and train the model.

```
[ ]: from sklearn.model_selection import train_test_split
     from sklearn.linear_model import LinearRegression

     # Define features (X) and target variable (y)
     X = df_encoded.drop(['Id', 'Price'], axis=1)
     y = df_encoded['Price']

     # Split the data into training and testing sets
     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     random_state=42)

     # Instantiate a LinearRegression model
     model = LinearRegression()

     # Train the LinearRegression model
     model.fit(X_train, y_train)
```

```
[ ]: LinearRegression()
```

1.5 Model evaluation

1.5.1 Subtask:

Evaluate the model's performance on the testing set using appropriate metrics like Mean Squared Error (MSE) and R-squared.

Reasoning: Import necessary metrics, make predictions, calculate and print MSE and R-squared, and interpret the results.

```
[ ]: from sklearn.metrics import mean_squared_error, r2_score

     # Make predictions on the test set
     y_pred = model.predict(X_test)

     # Calculate Mean Squared Error (MSE)
     mse = mean_squared_error(y_test, y_pred)
     print(f"Mean Squared Error (MSE): {mse}")
```

```

# Calculate R-squared score
r2 = r2_score(y_test, y_pred)
print(f"R-squared (R2): {r2}")

# Interpretation
print("\nModel Evaluation Interpretation:")
print(f"The Mean Squared Error (MSE) of {mse:.2f} indicates the average squared_
↪ difference between the actual and predicted prices. A lower MSE suggests_
↪ better model performance.")
print(f"The R-squared score of {r2:.2f} represents the proportion of the_
↪ variance in the house prices that is predictable from the features. An_
↪ R-squared value closer to 1.0 indicates that a larger proportion of the_
↪ variance is explained by the model, suggesting a better fit.")

```

Mean Squared Error (MSE): 78321466146.0328

R-squared (R2): -0.006717808430749761

Model Evaluation Interpretation:

The Mean Squared Error (MSE) of 78321466146.03 indicates the average squared difference between the actual and predicted prices. A lower MSE suggests better model performance.

The R-squared score of -0.01 represents the proportion of the variance in the house prices that is predictable from the features. An R-squared value closer to 1.0 indicates that a larger proportion of the variance is explained by the model, suggesting a better fit.